

Bank Customers Analysis



Table of contents

01 Introduction

Introduction of data and target

02 Data Analysis

Cleaning and Illustration

03 Prediction

Prediction by Machine Learning

04 Result

Summary





Introduction

Data Analysis

Prediction

Result

Introduction



Introduction



Dataset:

<https://www.kaggle.com/datasets/radheshyamkollipara/bank-customer-churn>

Tool:

- Python
- PBI
- Google slide

Target:

- Illustrate customers characters and suggest some business ideas accordingly.
- Predict customer trend.

Introduction

Data Analysis

Prediction

Result

Introduction

Shape: 10.000 rows & 17 columns

Customer Anthropologies (5 columns)

1. **CustomerId**
2. **Surname**
3. **Gender**
4. **Age**
5. **Geography**

Customer Characteristics (8 columns)

1. **Balance**
2. **EstimatedSalary**
3. **NumOfProducts**
4. **HasCrCard**
5. **IsActiveMember**
6. **Card Type**
7. **Tenure**

Customer Behaviors (4 columns)

1. **CreditScore**
2. **Complain**
3. **Satisfaction Score**
4. **Point Earned**
5. **Exited**

Introduction

```
cb.duplicated().sum()
```

✓ 0.0s

0

Data **does not have** duplicated rows or null values.

Checking columns individually is **clean**.

Moving on **Illustration**.

```
cb.isnull().sum()
```

✓ 0.0s

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0
Complain	0
Satisfaction Score	0
Card Type	0
Point Earned	0
dtype:	int64

```
cb.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore             10000 non-null  int64
4   Geography               10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                  10000 non-null  int64
8   Balance                 10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard               10000 non-null  int64
11  IsActiveMember          10000 non-null  int64
12  EstimatedSalary         10000 non-null  float64
13  Exited                  10000 non-null  int64
14  Complain                10000 non-null  int64
15  Satisfaction Score      10000 non-null  int64
16  Card Type               10000 non-null  object
17  Point Earned            10000 non-null  int64
dtypes: float64(2), int64(12), object(4)
memory usage: 1.4+ MB
```

Introduction

Data Analysis

Prediction

Result

Analysis



Introduction

Data Analysis


Prediction

Result

Overview

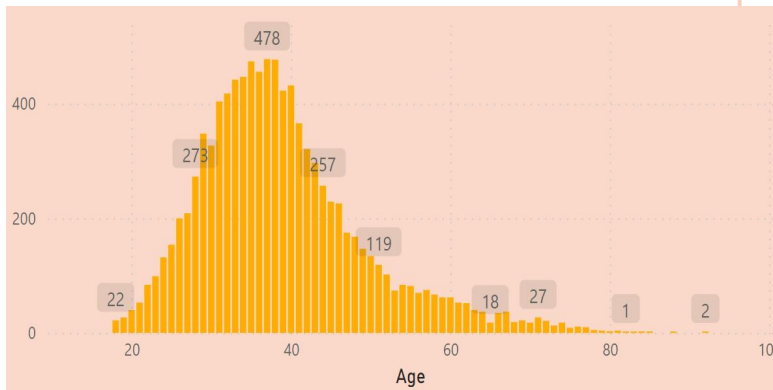
Dataset including
10.000 customers.

10.000
Customers



54% of them are
male.

Customers are mainly
from **30 to 45 years old.**



France

5.01K

Germany

2.51K

Spain

2.48K

Half of customers from
France. Others come
from **Germany** and
Spain.

Introduction

Data Analysis


Prediction

Result

Client comparision



Between current and lost clients.

Over 2.000 customers had left, it contains of **20%** percents of total customers.



Most customers are **Adults**. Number of **Senior loyal customers to their loss** is quite the **same**.

Banks have **more Males** customer than Females. But, **Females lost rate is higher**.



One thirds of German customer **left** the bank.

Introduction

Data Analysis

Prediction

Result


Client comparision

Between current and lost clients.

The fewer the products **higher** the customers **churn** customers.




Number of customer remains **uniform** across all categories.



Average **balance** of **Loyal** customer is **lower** than Lost customer.



Not active customer are more likely to leave.



Introduction

Data Analysis


Prediction

Result

Client comparision

Between current and lost clients.

Numbers of customers in each satisfaction score are **similar**.



Customer using **Credit Card** has **higher leaving rate** than customer who not own credit card.

Lost clients has much **lower salary range** than current clients.



Credit score range of **600 - 700** has the **highest current customers** also has **highest lost customers** rate.

- Customers **mostly adult** from 30 to 45 years old, **churn rate** of this adult also **really high of 66%**.
- **Half** of the customer are **Frances**. German account of 25% however their **leaving rate nearly 40%**.
- There are **more male than female** but **leaving rate of female is higher**.
- Most customers use 1 to 2 products. **Churn rate of 1 product user is higher than the rest**.
- The amount of customers in each **Card Types pretty corresponding**, Gold customers slightly higher.
- Average **balance** of Loyal customers much **lower** than Lost customers, but their **salary is higher**.
- Customers **using Credit Card** have **higher leaving rate** than customers who not own credit card.
Credit score from 600 to 700 account 30% of customers highest lost customers rate of 7%.
- **Satisfaction score does not effect** bank churn rate.

Introduction

Data Analysis

Prediction

Result

Prediction

Preparation



Encoding

encoding

```
for i in cb.columns:
    print(i)
    print(cb[i].unique())
    print('-----')

# there are 3 columns need to encode

# 'Geography': 3 unique
# 'Gender' : 2 unique
# 'Card Type' : 4 unique
```

✓ 0.0s

Outputs are collapsed ...

```
# encode 3 columns have string unique:

dictionary_gender = {'Male' : 0, 'Female' : 1}
cb['Gender_encode'] = cb.Gender.map(dictionary_gender)

onehot_Geography = pd.get_dummies(cb['Geography'], prefix = 'Geography').astype(int)
onehot_Card_type = pd.get_dummies(cb['Card Type'], prefix = 'Card Type').astype(int)
```

✓ 0.0s


There are **3 columns** have string values and need to encode.

- Column '**Gender**' applied map dictionary method.
- Columns '**Geography**' and '**Card_type**' applied onehot method.

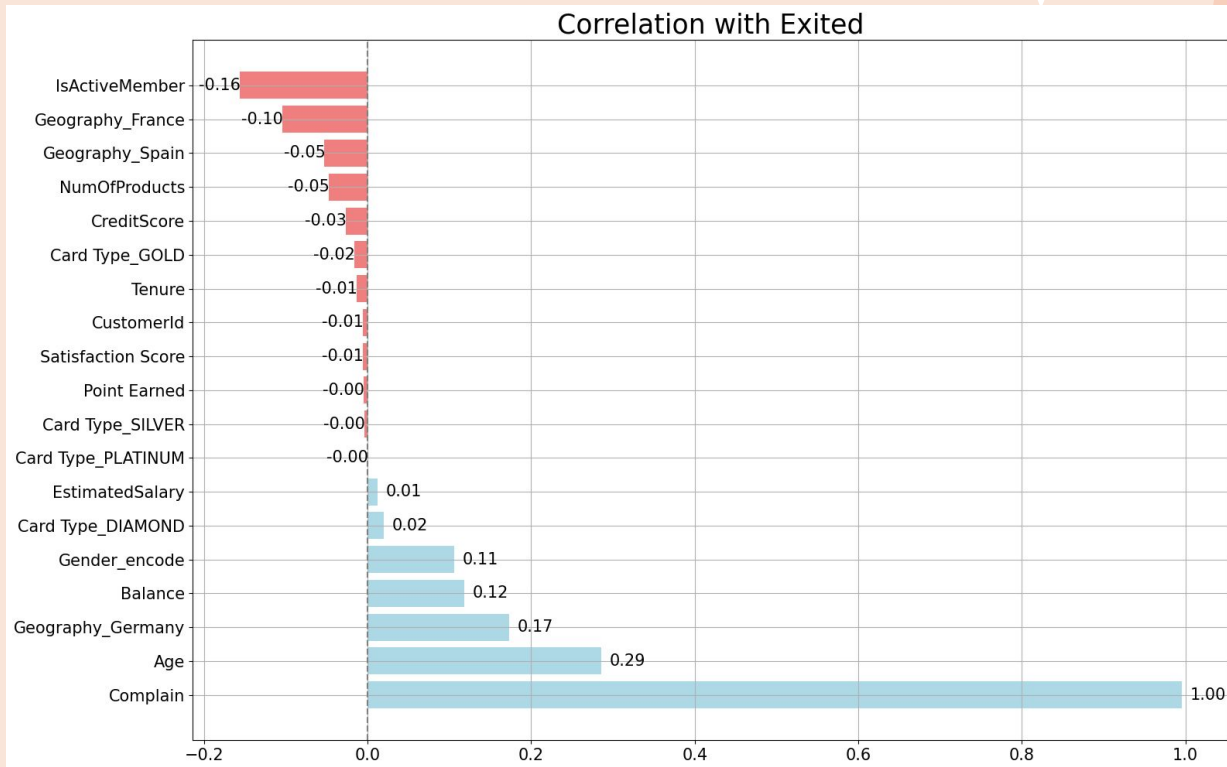
X, y defining

Complain has absolute correlation.
Customers who make complain more likely leave the bank. Current Clients only make 10 complains/ 7962 clients. In contrast, 2034 complains/ 2038 clients who left the bank. -> Not use.

● Complain ● No comment



As showed, correlation rate around -0.16 to 0.29.




X, y defining

Checking correlation:

After **filter out** all of the columns have **high correlation** columns with columns '**Exited**' which is **below -0.1** and **higher than 0.1**, then **check** their **correlation with each other**.

The result looks good. Continue.



★ X, y defining


Select X, y for model:

```
# chose x, y
y = cbc['Exited'].values
X = cbc[['Age', 'Balance', 'Gender_encode', 'Geography_France', 'IsActiveMember']].values
```

✓ 0.0s

Python

★ Imbalance data



For the **data balance** of columns '**Exited**', the **0 values** (current clients) are much higher than **1 values** (exited clients).

Choose the **undersize method** to process.

Normalization

normalization: min-max scaler

```
# Scale X2
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(cb_balance.iloc[:, 1:11].values)

# create new dataframe
cbc = pd.DataFrame(data = X_scaled, columns = cb_balance.iloc[:, 1:11].columns)

# add y2 column
cbc['Exited'] = cb_balance['Exited']

cbc.head()
```

✓ 0.0s

	Age	Balance	Gender_encode	Geography_France	Geography_Germany	IsActiveMember	Exited
0	0.200	0.000	1.000	1.000	0.000	1.000	0
1	0.286	0.000	0.000	1.000	0.000	0.000	0
2	0.157	0.000	0.000	0.000	0.000	0.000	0
3	0.257	0.445	0.000	0.000	0.000	1.000	0
4	0.371	0.456	1.000	1.000	0.000	1.000	0

Split Train-Test

split train test dataset

```
# chose x, y
y = cbc['Exited'].values
x = cbc[['Age', 'Balance', 'Gender_encode', 'Geography_France', 'Geography_Germany', 'IsActiveMember']].values
```

✓ 0.0s

```
# create X_set for later use in model result
X_set = ['Age', 'Balance', 'Gender_encode', 'Geography_France', 'Geography_Germany', 'IsActiveMember']
```

✓ 0.0s

```
# run split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
```

✓ 0.0s

```
print(f'Data X_train: {X_train.sum()}')
print(f'Data y_train: {y_train.sum()}')
print(f'Data X_test: {X_test.sum()}')
print(f'Data y_test: {y_test.sum()}')
```

✓ 0.0s

```
Data X_train: 6768.618920505589
Data y_train: 1438
Data X_test: 2882.1677445508994
Data y_test: 600
```

Introduction

Data Analysis

Prediction

Result

Prediction

Machine learning




Models:

Models	Parameter
Logistic Regression	Default
Gaussian Naive Bayes	Default
Decision Tree	max_depth = 6
Random Forest	n_estimators = 71
K Nearst Neighbor	n_neighbors = 30


Logistic Regression

	precision	recall	f1-score	support
0	0.73	0.72	0.72	623
1	0.71	0.72	0.71	600
accuracy			0.72	1223
macro avg	0.72	0.72	0.72	1223
weighted avg	0.72	0.72	0.72	1223





Gaussian Naive Bayes

	precision	recall	f1-score	support
0	0.69	0.70	0.69	623
1	0.68	0.67	0.67	600
accuracy			0.68	1223
macro avg	0.68	0.68	0.68	1223
weighted avg	0.68	0.68	0.68	1223

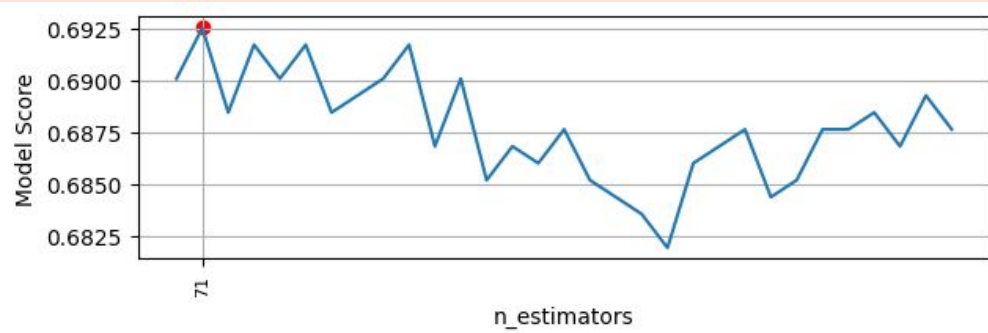


Decision Tree



	precision	recall	f1-score	support
0	0.73	0.74	0.73	623
1	0.73	0.71	0.72	600
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

Random Forest



	precision	recall	f1-score	support
0	0.6976	0.6998	0.6987	623
1	0.6873	0.6850	0.6861	600
accuracy			0.6926	1223
macro avg	0.6924	0.6924	0.6924	1223
weighted avg	0.6925	0.6926	0.6925	1223


Confusion Matrix - Forest model

True label	Predicted label	
	Clients	Exited
Clients	436	187
Exited	189	411

K Nearest Neighbors



	precision	recall	f1-score	support
0	0.7250	0.7785	0.7508	623
1	0.7509	0.6933	0.7210	600
accuracy			0.7367	1223
macro avg	0.7379	0.7359	0.7359	1223
weighted avg	0.7377	0.7367	0.7362	1223



Introduction

Data Analysis

Prediction

Result

Result



Introduction

Data Analysis

Prediction


Result

Model	Accuracy Score	F1_Score	Precision	Recall	ROC AUC Score
Random forest	0.69	0.69	0.69	0.69	0.76
Logistic Regression	0.72	0.71	0.71	0.72	0.78
KNN	0.74	0.72	0.75	0.69	0.80
GNB	0.68	0.67	0.68	0.67	0.76
Decesion Tree	0.73	0.72	0.73	0.71	0.78

KNN have the **highest result** in all score.

Decesion Tree is the **2nd best performance**.

Recomended this 2 models in this dataset.



Thank you for reading!

01 Introduction

Introduction of data and target

02 Data Analysis

Cleaning and Illustration

03 Prediction

Prediction by Machine Learning

04 Result

Summary

